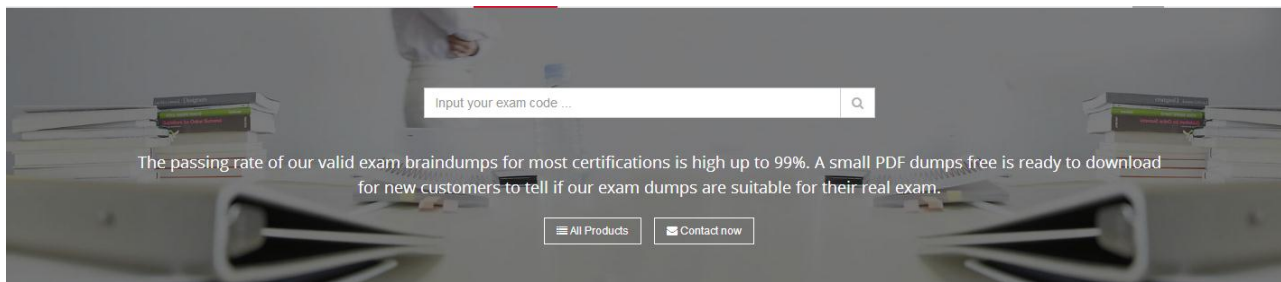


# RealExamFree



## Why Choose Us



### QUALITY AND VALUE

RealExamFree Practice Exams are written to the highest standards of technical accuracy, using only certified subject matter experts and published authors for development - no all dumps.



### TESTED AND APPROVED

We are committed to the process of vendor and third party approvals. We believe professionals and executives alike deserve the confidence of quality coverage these authorizations provide.



### EASY TO PASS

If you prepare for the exams using our RealExamFree testing engine, It is easy to succeed for all certifications in the first attempt. You don't have to deal with all dumps or any free torrent / rapidshare all stuff.



### TRY BEFORE BUY

RealExamFree offers free demo of each product. You can check out the interface, question quality and usability of our practice exams before you decide to buy.

## Try Before You Buy

Download a free sample of any of our exam questions and answers

- ✓ 24/7 customer support, Secure shopping site
- ✓ Free One year updates to match real exam scenarios
- ✓ If you failed your exam after buying our products we will refund the full amount back to you.

[Free Download](#)

## Customer Reviews



The exam is easy. many questions are same with practice paper before. Pass it easily

**Dominic**



It is worthy it. I am happy about my score. Thank you for the dumps.

**Gloria**



Don't waste too much time on what you are not good at. Let others help you. Yes. I am lucky to order this exam cram and pass test casually. Wonderful

**Nat**



It is very a good dumps. It is same with real exam. best dumps without one of. the only.

**Joanna**

<http://www.realexamfree.com>

Most international real exam & valid exam braindumps & dumps free

**Exam** : **DP-200**

**Title** : Implementing an Azure Data Solution

**Vendor** : Microsoft

**Version** : DEMO

**NO.1** You need to develop a pipeline for processing data. The pipeline must meet the following requirements.

- \* Scale up and down resources for cost reduction.
- \* Use an in-memory data processing engine to speed up ETL and machine learning operations.
- \* Use streaming capabilities.
- \* Provide the ability to code in SQL, Python, Scala, and R.
- \* Integrate workspace collaboration with Git.

What should you use?

- A.** HDInsight Spark Cluster
- B.** Azure SQL Data Warehouse
- C.** HDInsight Hadoop Cluster
- D.** Azure Stream Analytics

**Answer:** A

Explanation:

Apache Spark is an open-source, parallel-processing framework that supports in-memory processing to boost the performance of big-data analysis applications.

HDInsight is a managed Hadoop service. Use it to deploy and manage Hadoop clusters in Azure. For batch processing, you can use Spark, Hive, Hive LLAP, MapReduce.

Languages: R, Python, Java, Scala, SQL

You can create an HDInsight Spark cluster using an Azure Resource Manager template. The template can be found in GitHub.

References:

<https://docs.microsoft.com/en-us/azure/architecture/data-guide/technology-choices/batch-processing>

**NO.2** Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution. After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You need to implement diagnostic logging for Data Warehouse monitoring.

Which log should you use?

- A.** DmsWorkers
- B.** SqlRequests
- C.** RequestSteps
- D.** ExecRequests

**Answer:** B

Explanation:

Scenario:

The Azure SQL Data Warehouse cache must be monitored when the database is being used.

Metric	Description
A	Low cache hit %, high cache usage %
B	Low cache hit %, low cache usage %
C	High cache hit %, high cache usage %

References:

<https://docs.microsoft.com/en-us/sql/relational-databases/system-dynamic-management-views/sys-dm-pdw-sql-requests-transact-sq>

**NO.3** You are designing a new Lambda architecture on Microsoft Azure.

The real-time processing layer must meet the following requirements:

Ingestion:

- \* Receive millions of events per second
- \* Act as a fully managed Platform-as-a-Service (PaaS) solution
- \* Integrate with Azure Functions

Stream processing:

- \* Process on a per-job basis
- \* Provide seamless connectivity with Azure services
- \* Use a SQL-based query language

Analytical data store:

- \* Act as a managed service
- \* Use a document store
- \* Provide data encryption at rest

You need to identify the correct technologies to build the Lambda architecture using minimal effort.

Which technologies should you use? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

### Architecture requirement

### Answer Area

#### Ingestion

HDInsight Kafka	V
Azure Event Hubs	
HDInsight Storm	
HDInsight Spark	

#### Stream Processing

Azure Stream Analytics	V
HDInsight with Spark Streaming	
Azure Cosmos DB Change Feed	
Azure Analysis Services	

#### Analytical Data Store

Hive LLAP on HDInsight	V
Azure Analysis Services	
Azure Cosmos DB	
SQL Data Warehouse	

**Answer:****Architecture requirement****Answer Area****Ingestion**

HDInsight Kafka	✓
Azure Event Hubs	
HDInsight Storm	
HDInsight Spark	

**Stream Processing**

Azure Stream Analytics	✓
HDInsight with Spark Streaming	
Azure Cosmos DB Change Feed	
Azure Analysis Services	

**Analytical Data Store**

Hive LLAP on HDInsight	✓
Azure Analysis Services	
Azure Cosmos DB	
SQL Data Warehouse	

Explanation:

Box 1: Azure Event Hubs

This portion of a streaming architecture is often referred to as stream buffering. Options include Azure Event Hubs, Azure IoT Hub, and Kafka.

Incorrect Answers: Not HDInsight Kafka

Azure Functions need a trigger defined in order to run. There is a limited set of supported trigger types, and Kafka is not one of them.

Box 2: Azure Stream Analytics

Azure Stream Analytics provides a managed stream processing service based on perpetually running SQL queries that operate on unbounded streams.

You can also use open source Apache streaming technologies like Storm and Spark Streaming in an HDInsight cluster.

Box 3: Azure SQL Data Warehouse

Azure SQL Data Warehouse provides a managed service for large-scale, cloud-based data warehousing. HDInsight supports Interactive Hive, HBase, and Spark SQL, which can also be used to serve data for analysis.

References:

<https://docs.microsoft.com/en-us/azure/architecture/data-guide/big-data/>

**NO.4** Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution. Determine whether the solution meets the stated goals.

You develop a data ingestion process that will import data to a Microsoft Azure SQL Data Warehouse. The data to be ingested resides in parquet files stored in an Azure Data Lake Gen 2 storage account. You need to load the data from the Azure Data Lake Gen 2 storage account into the Azure SQL Data Warehouse.

Solution:

1. Create an external data source pointing to the Azure storage account
2. Create an external file format and external table using the external data source
3. Load the data using the INSERT...SELECT statement

Does the solution meet the goal?

- A. No
- B. Yes

**Answer:** A

Explanation:

You load the data using the CREATE TABLE AS SELECT statement.

References:

<https://docs.microsoft.com/en-us/azure/sql-data-warehouse/sql-data-warehouse-load-from-azure-data-lake-store>

**NO.5** A company runs Microsoft SQL Server in an on-premises virtual machine (VM).

You must migrate the database to Azure SQL Database. You synchronize users from Active Directory to Azure Active Directory (Azure AD).

You need to configure Azure SQL Database to use an Azure AD user as administrator. What should you configure?

- A. For each Azure SQL Database, set the Access Control to administrator.
- B. For the Azure SQL Database server, set the Active Directory to administrator.
- C. For each Azure SQL Database, set the Active Directory administrator role.
- D. For the Azure SQL Database server, set the Access Control to administrator.

**Answer:** C

Explanation:

There are two administrative accounts (Server admin and Active Directory admin) that act as administrators.

One Azure Active Directory account, either an individual or security group account, can also be configured as an administrator. It is optional to configure an Azure AD administrator, but an Azure AD administrator must be configured if you want to use Azure AD accounts to connect to SQL Database.

References:

<https://docs.microsoft.com/en-us/azure/sql-database/sql-database-manage-logins>

**NO.6** You develop data engineering solutions for a company.

A project requires the deployment of data to Azure Data Lake Storage.

You need to implement role-based access control (RBAC) so that project members can manage the Azure Data Lake Storage resources.

Which three actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Configure access control lists (ACL) for the Azure Data Lake Storage account.
- B. Configure service-to-service authentication for the Azure Data Lake Storage account.
- C. Assign Azure AD security groups to Azure Data Lake Storage.
- D. Configure end-user authentication for the Azure Data Lake Storage account.
- E. Create security groups in Azure Active Directory (Azure AD) and add project members.

**Answer:** A,C,E

Explanation:

References:

<https://docs.microsoft.com/en-us/azure/data-lake-store/data-lake-store-secure-data>

**NO.7** You are developing a solution using a Lambda architecture on Microsoft Azure.

The data at test layer must meet the following requirements:

Data storage:

- \* Serve as a repository (or high volumes of large files in various formats).
- \* Implement optimized storage for big data analytics workloads.
- \* Ensure that data can be organized using a hierarchical structure.

Batch processing:

- \* Use a managed solution for in-memory computation processing.
- \* Natively support Scala, Python, and R programming languages.
- \* Provide the ability to resize and terminate the cluster automatically.

Analytical data store:

- \* Support parallel processing.
- \* Use columnar storage.
- \* Support SQL-based languages.

You need to identify the correct technologies to build the Lambda architecture.

Which technologies should you use? To answer, select the appropriate options in the answer area

NOTE: Each correct selection is worth one point.



**Architecture requirement****Technology**

Data storage

	▼
Azure SQL Database	
Azure Blob Storage	
Azure Cosmos DB	
Azure Data Lake Store	

Batch processing

	▼
HDInsight Spark	
HDInsight Hadoop	
Azure Databricks	
HDInsight Interactive Query	

Analytical data store

	▼
HDInsight HBase	
Azure SQL Data Warehouse	
Azure Analysis Services	
Azure Cosmos DB	

**Answer:**



**Architecture requirement****Technology**

Data storage

	▼
Azure SQL Database	
Azure Blob Storage	
Azure Cosmos DB	
Azure Data Lake Store	

Batch processing

	▼
HDInsight Spark	
HDInsight Hadoop	
Azure Databricks	
HDInsight Interactive Query	

Analytical data store

	▼
HDInsight HBase	
Azure SQL Data Warehouse	
Azure Analysis Services	
Azure Cosmos DB	

Explanation:

Data storage: Azure Data Lake Store

A key mechanism that allows Azure Data Lake Storage Gen2 to provide file system performance at object storage scale and prices is the addition of a hierarchical namespace. This allows the collection of objects/files within an account to be organized into a hierarchy of directories and nested subdirectories in the same way that the file system on your computer is organized. With the hierarchical namespace enabled, a storage account becomes capable of providing the scalability and cost-effectiveness of object storage, with file system semantics that are familiar to analytics engines and frameworks.

Batch processing: HD Insight Spark

Apache Spark is an open-source, parallel-processing framework that supports in-memory processing to boost the performance of big-data analysis applications.

HDInsight is a managed Hadoop service. Use it to deploy and manage Hadoop clusters in Azure. For batch processing, you can use Spark, Hive, Hive LLAP, MapReduce.

Languages: R, Python, Java, Scala, SQL

Analytic data store: SQL Data Warehouse

SQL Data Warehouse is a cloud-based Enterprise Data Warehouse (EDW) that uses Massively Parallel Processing (MPP).

SQL Data Warehouse stores data into relational tables with columnar storage.

References:

<https://docs.microsoft.com/en-us/azure/storage/blobs/data-lake-storage-namespace>

<https://docs.microsoft.com/en-us/azure/architecture/data-guide/technology-choices/batch-processing>

<https://docs.microsoft.com/en-us/azure/sql-data-warehouse/sql-data-warehouse-overview-what-is>

**NO.8** You manage a solution that uses Azure HDInsight clusters.

You need to implement a solution to monitor cluster performance and status.

Which technology should you use?

- A.** Azure Log Analytics
- B.** Ambari REST API
- C.** Azure HDInsight REST API
- D.** Ambari Web UI
- E.** Azure HDInsight .NET SDK

**Answer:** D

Explanation:

Ambari is the recommended tool for monitoring utilization across the whole cluster. The Ambari dashboard shows easily glanceable widgets that display metrics such as CPU, network, YARN memory, and HDFS disk usage. The specific metrics shown depend on cluster type. The "Hosts" tab shows metrics for individual nodes so you can ensure the load on your cluster is evenly distributed. The Apache Ambari project is aimed at making Hadoop management simpler by developing software for provisioning, managing, and monitoring Apache Hadoop clusters. Ambari provides an intuitive, easy-to-use Hadoop management web UI backed by its RESTful APIs.

References:

<https://azure.microsoft.com/en-us/blog/monitoring-on-hdinsight-part-1-an-overview/>

<https://ambari.apache.org/>